

EE456 Section 001

Mini Project 2

Report

CNN TO CLASSIFY CIFAR10 DATASET

Sahin, Efe

PENN STATE UNIVERSITY | DECEMBER 5, 2022

Overview

This Project is about using CNN to solve the image classification problem where the CIFAR10 dataset will be used. CIFAR10 is an image dataset of 32x32 images with 3 channels (R,G,B). There are in total 60000 images which consist of 10000 per class (therefore there are 10 classes). Each class represents a specific category like, bird, truck, ship etc. The CNN is trained based on a training set randomly augmented (for this project augmentations include random flips and random crops) and selected from the main dataset.

Implementation

It is important to mention that the code I ran “cifar10NN.vfinal.ipynb” uses python 3.8.15 with several other libraries like torch, torchvision, matplotlib and numpy inside of an Anaconda environment. I have installed the following versions to benefit the GPU of my Mac M1 PC.

- Torch: 1.14.0.dev20221130
- Torchvision: 0.15.0.dev20221130
- Matplotlib: (latest is fine)
- Numpy: 1.24.0rc1 (latest is fine)

The code is able to select the proper device to maximize the possible speed benefits.

```
1 if torch.backends.mps.is_available:
2     processor='mps'
3 elif torch.cuda.is_available():
4     processor='cuda0'
5 else:
6     processor='cpu'
7 device = torch.device(device=processor)
8 print(f"device is set as: {device}")
```

✓ 0.1s

device is set as: mps

FIGURE 1: MPS IS SPECIAL TO M1 APPLE COMPUTERS, CUDA:0 IS THE OTHER OPTION, CPU IS THE LAST OPTION.

Then it is time to split the main dataset into train, validation, testing sets. I decided to use 50000 randomly picked and augmented images from the main dataset then split it into 35000 and 15000 images of train and validation data sets respectively. The remaining 10000 images will be used for testing after the Net has been trained.

```

1  normalized = transforms.Normalize((0.49139968, 0.48215827, 0.44653124), (0.24703233, 0.24348505, 0.26158768))
2  flip = transforms.RandomHorizontalFlip()
3  crop = transforms.RandomCrop(size=32)
4
5  trainValidTransform = transforms.Compose([transforms.ToTensor(), normalized, flip, crop])
6  testTransform = transforms.Compose([transforms.ToTensor(), normalized])
7
8  trainValidSet = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=trainValidTransform)
9  trainSet, validSet = data.random_split(trainValidSet, [35000, 15000])
10 testSet = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=testTransform)
11
12 trainLoader = torch.utils.data.DataLoader(trainSet, batch_size=batchSize, shuffle=True, num_workers=2)
13 validLoader = torch.utils.data.DataLoader(validSet, batch_size=batchSize, shuffle=True, num_workers=2)
14 testLoader = torch.utils.data.DataLoader(testSet, batch_size=batchSize, shuffle=False, num_workers=2)
15
16 classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
17 print('Train data set:', len(trainSet))
18 print('Valid data set:', len(validSet))
19 print('Test data set:', len(testSet))

```

✓ 1.5s

Files already downloaded and verified
Files already downloaded and verified
Train data set: 35000
Valid data set: 15000
Test data set: 10000

FIGURE 2: FIRST NORMALIZATION OF CIFAR10, FLIP, CROP ARE DECIDED IN LINES 1-3. THEN TRAIN, VALID, TESTING SETS HAVE BEEN FORMED IN LINES 5-10. THEN THEY ARE LOADED IN LINES 12-14.

The Structure of the neural net consists of 3 convolutional layers, 2 pooling and followed by 3 fully connected linear layers. The Net is defined as the class ConvNet1 and made as the model object. Pytorch can show detailed summary for cuda0 or cpu processors but not for mps. Then for loss, CrossEntropyLoss is used since this net will classify multiple classes. As optimizer SGD i.e stochastic gradient descent is used to optimize the parameters of the net.




FIGURE 3: A CNN HEURISTIC.

```

1  import torch.nn as nn
2  import torch.nn.functional as F
3
4  class ConvNet1(nn.Module):
5      def __init__(self):
6          super().__init__()
7          self.conv1 = nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, stride=1, padding=1) # 32,32,32
8          self.conv2 = nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, stride=1, padding=1) # 64,32,32
9          self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2, padding=0) # 64,16,16
10         self.conv3 = nn.Conv2d(in_channels=64, out_channels=94, kernel_size=3, stride=1, padding=1) # 94,16,16
11         self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2, padding=0) # 94,8,8
12
13         self.fc1 = nn.Linear(in_features=94*8*8, out_features=128)
14         self.fc2 = nn.Linear(in_features=128, out_features=64)
15         self.fc3 = nn.Linear(in_features=64, out_features=10)
16
17     def forward(self, x):
18         x = F.relu(self.conv1(x))
19         x = F.relu(self.conv2(x))
20         x = self.pool1(x)
21         x = F.relu(self.conv3(x))
22         x = self.pool2(x)
23
24         x = torch.flatten(x, start_dim=1)
25         x = F.relu(self.fc1(x))
26         x = F.relu(self.fc2(x))
27         x = self.fc3(x)
28         return x
29
30 model = ConvNet1().to(device)
31
32 if processor=='cpu' or processor=='cuda0':
33     from torchsummary import summary
34     summary(model, (3, 32, 32))
35
36 criterion = nn.CrossEntropyLoss()
37 optimizer = optim.SGD(model.parameters(), lr=learnRate, momentum=0.9)

```

FIGURE 4: THE CLASS OF COVNET1.

Results/Conclusion

During the training process of the CNN for 15 epochs with 20 image batches, I displayed the trainingloss, validationloss, trainingaccuracy and validationaccuracies. It seemed from the values that 15 epochs were enough to achieve training accuracy above 90. I also realized that training accuracy surpassed validation accuracy. Similar relationship can be seen from the calculated loss values.

For epoch: 1	---	trainloss=1.653847	validloss=1.238903		trainaccuracy=39.545715	valid=55.360001
For epoch: 2	---	trainloss=1.102419	validloss=1.069913		trainaccuracy=60.237144	valid=63.366669
For epoch: 3	---	trainloss=0.882153	validloss=0.813558		trainaccuracy=68.842857	valid=71.366669
For epoch: 4	---	trainloss=0.755343	validloss=0.804646		trainaccuracy=73.371429	valid=72.546669
For epoch: 5	---	trainloss=0.655973	validloss=0.813546		trainaccuracy=77.405716	valid=72.139999
For epoch: 6	---	trainloss=0.579830	validloss=0.780484		trainaccuracy=79.902855	valid=73.986664
For epoch: 7	---	trainloss=0.497934	validloss=0.805496		trainaccuracy=82.534286	valid=74.146667
For epoch: 8	---	trainloss=0.445979	validloss=0.735752		trainaccuracy=84.494286	valid=76.153336
For epoch: 9	---	trainloss=0.392063	validloss=0.876399		trainaccuracy=86.491432	valid=74.366669
For epoch: 10	---	trainloss=0.353841	validloss=0.790553		trainaccuracy=88.011429	valid=75.453331
For epoch: 11	---	trainloss=0.315596	validloss=0.879806		trainaccuracy=89.245712	valid=74.693336
For epoch: 12	---	trainloss=0.268274	validloss=0.910542		trainaccuracy=90.785713	valid=75.566666
For epoch: 13	---	trainloss=0.256548	validloss=1.030078		trainaccuracy=91.314285	valid=74.546669
For epoch: 14	---	trainloss=0.224269	validloss=1.036487		trainaccuracy=92.540001	valid=75.040001
For epoch: 15	---	trainloss=0.212034	validloss=1.004622		trainaccuracy=92.868568	valid=75.306671
Finished Training						

FIGURE 5: EPOCHS AND RESPECTIVE STATISTICS OF THE CNN DURING TRAINING

The following figures 6 and 7 show the plots of loss and accuracy for the CNN training.




FIGURE 6




FIGURE 7

Later the CNN is put into the testing process with the testing dataset. The following accuracy of 75% has been observed, which is very similar accuracy percentage from the validation. The figure 8 also shows the accuracy levels of predicting each class correctly

```
Accuracy of the network on the 10000 test images: 75.1 %
Accuracy for class: plane is 80.4 %
Accuracy for class: car is 87.9 %
Accuracy for class: bird is 59.3 %
Accuracy for class: cat is 58.9 %
Accuracy for class: deer is 74.6 %
Accuracy for class: dog is 53.9 %
Accuracy for class: frog is 83.1 %
Accuracy for class: horse is 81.2 %
Accuracy for class: ship is 88.2 %
Accuracy for class: truck is 83.5 %
```

FIGURE 8:

Then I selected a random batch of 20 images to display with their labels vs what the neural net predicted.

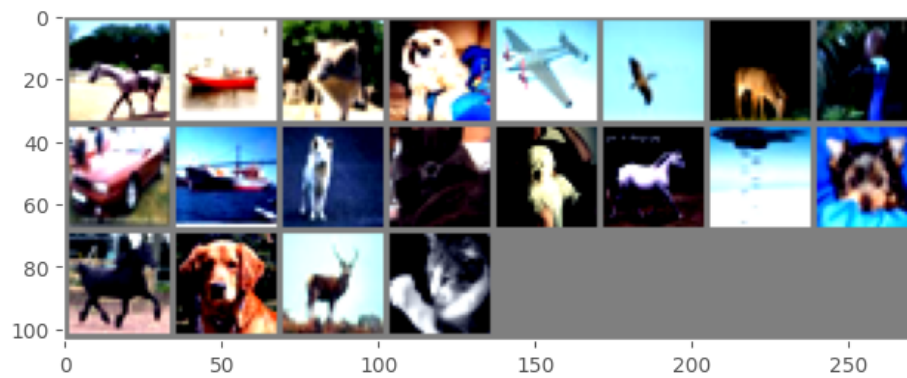


FIGURE 9: A BATCH OF IMAGES

Actualval:	horse	ship	bird	dog	plane	bird	deer	bird	car	ship	dog	cat	dog	horse	plane	dog
	horse	dog	deer	cat												
Predicted:	horse	ship	bird	dog	plane	plane	deer	bird	car	ship	horse	cat	dog	horse	plane	dog
	horse	dog	deer	cat												

FIGURE 10: THE BATCH OF IMAGES WITH THEIR RESPECTIVE ACTUAL LABELS AND PREDICTED VALUES.

Lastly the confusion matrix resulted a plot of Actual labels of classes vs the number of predictions based on the training set. It seems that the diagonally, meaning the correct predictions, have been the majority as those boxes were darker than others.




FIGURE 11: CONFUSION MATRIX OF CIFAR10 ACCORDING TO THE CNN TRAINED IN THIS PROJECT